```
from transformers import pipeline, set_seed, GPT2LMHeadModel, GPT2Tokenizer
import torch

# Reproducibility for exact generations
set_seed(42)

generator = pipeline('text-generation', model='gpt2',
                     device=0 if torch.cuda.is_available() else -1)

prompt = "Once upon a time in a distant galaxy,"
outputs = generator(prompt, max_new_tokens=100, num_return_sequences=3,
                    temperature=0.8, top_p=0.95, do_sample=True, pad_token_id=50256)

print("=== Text Generations from Pipeline ===")
for i, output in enumerate(outputs):
    print(f"\nGeneration {i+1}:\n{output['generated_text']}")
    print("-" * 80)

# Manual approach with more control
tokenizer = GPT2Tokenizer.from_pretrained('gpt2')
model = GPT2LMHeadModel.from_pretrained('gpt2')

# Move the model to the correct device to avoid RuntimeError
model.to(generator.device)

tokenizer.pad_token = tokenizer.eos_token

prompt2 = "In a world where AI has become conscious,"
inputs = tokenizer(prompt2, return_tensors="pt").to(generator.device)

outputs2 = model.generate(**inputs, max_new_tokens=120,
                          temperature=0.7, top_p=0.9, do_sample=True, repetition_penalty=1.2)

print("\n=== Manual Generation with Full Control ===")
print(tokenizer.decode(outputs2[0], skip_special_tokens=True))
```

```
Loading weights: 100%|| 148/148 [00:00<00:00, 490.61it/s, Materializing param=transformer.wte.weight]
GPT2LMHeadModel LOAD REPORT from: gpt2
Key                  | Status     |  |
---------------------+------------+--+-
h.{0...11}.attn.bias | UNEXPECTED |  |

Notes:
- UNEXPECTED    :can be ignored when loading from different task/architecture; not ok if you
Both `max_new_tokens` (=100) and `max_length`(=50) seem to have been set. `max_new_tokens` wi
=== Text Generations from Pipeline ===

Generation 1:
Once upon a time in a distant galaxy, a man in his thirties, in his early thirties, was broug
--------------------------------------------------------------------------------

Generation 2:
Once upon a time in a distant galaxy, there was nothing but the light of a light bulb and the

It is a small light bulb, and a small light bulb is very small.

The smallest light bulb of any kind was the most valuable part of the universe.

The most valuable part of the universe was what the universe made of light.

The most valuable
--------------------------------------------------------------------------------

Generation 3:
Once upon a time in a distant galaxy, the galaxy will pass through a new stage where life wil

This could have been a huge step forward for our understanding of how our universe works. It
--------------------------------------------------------------------------------
Loading weights: 100%|| 148/148 [00:00<00:00, 458.54it/s, Materializing param=transformer.wte.weight]
GPT2LMHeadModel LOAD REPORT from: gpt2
Key                  | Status     |  |
---------------------+------------+--+-
h.{0...11}.attn.bias | UNEXPECTED |  |

Notes:
- UNEXPECTED    :can be ignored when loading from different task/architecture; not ok if you
Setting `pad_token_id` to `eos_token_id`:50256 for open-end generation.

=== Manual Generation with Full Control ===
In a world where AI has become conscious, we can only imagine how much of an impact it will h
The more I thought about this subject in my mind the clearer and clear what happened to me as
```